

Full 500×4096^2 specimen, NVIDIA L40 — alternating before/after pairs
Shared GPU: observed timings, not an isolated end-to-end speedup claim

