



Production-scale before/after — three complete GPU runs per version, seeds 103–105
NVIDIA L40; black dots are individual runs; driver samples include CuPy and CUDA overhead

Complete CLI (warm caches)



Peak PyTorch GPU allocation



Sampled process GPU memory

